

From workflows to pipelines

DESIGNING MACHINE LEARNING WORKFLOWS IN PYTHON



Dr. Chris Anagnostopoulos
Honorary Associate Professor

Revisiting our workflow

```
from sklearn.ensemble import RandomForestClassifier as rf
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

```
grid_search = GridSearchCV(rf(), param_grid={'max_depth': [2, 5, 10]})
grid_search.fit(X_train, y_train)
depth = grid_search.best_params_['max_depth']
```

```
vt = SelectKBest(f_classif, k=3).fit(X_train, y_train)
clf = rf(max_depth=best_value).fit(vt.transform(X_train), y_train)
accuracy_score(clf.predict(vt.transform(X_test)), y_test)
```

The power of grid search

Optimize `max_depth` :

```
pg = {'max_depth': [2, 5, 10]}
gs = GridSearchCV(rf(),
                  param_grid=pg)
gs.fit(X_train, y_train)
depth = gs.best_params_['max_depth']
```

	10 estimators (default)	20 estimators	30 estimators
depth 2		?	?
depth 5		?	?
depth 10	★	?	?

The power of grid search

Then optimize `n_estimators` :

```
pg = {'n_estimators': [10, 20, 30]}
gs = GridSearchCV(
    rf(max_depth=depth),
    param_grid=pg)
gs.fit(X_train, y_train)
n_est = gs.best_params_[
    'n_estimators']
```

	10 estimators (default)	20 estimators	30 estimators
depth 2		?	?
depth 5		?	?
depth 10			★


The power of grid search

Jointly `max_depth` and `n_estimators` :

```
pg = {  
    'max_depth': [2, 5, 10],  
    'n_estimators': [10, 20, 30]  
}  
  
gs = GridSearchCV(rf(),  
    param_grid=pg)  
gs.fit(X_train, y_train)  
print(gs.best_params_)  
  
{ 'max_depth': 10, 'n_estimators': 20 }
```

	10 estimators (default)	20 estimators	30 estimators
depth 2			
depth 5			
depth 10			★

Pipelines



Pipelines

`Pipeline(n_estimators, max_depth, k)`

Random Forest
- `n_estimators`
- `max_depth`



Select K Best
- `k`

Pipelines

```
from sklearn.pipeline import Pipeline
pipe = Pipeline([
    ('feature_selection', SelectKBest(f_classif)),
    ('classifier', RandomForestClassifier())
])

params = dict(
    feature_selection__k=[2, 3, 4],
    classifier__max_depth=[5, 10, 20]
)

grid_search = GridSearchCV(pipe, param_grid=params)
gs = grid_search.fit(X_train, y_train).best_params_
```

```
{'classifier__max_depth': 20, 'feature_selection__k': 4}
```


Customizing your pipeline

```
from sklearn.metrics import roc_auc_score, make_scorer  
auc_scorer = make_scorer(roc_auc_score)  
grid_search = GridSearchCV(pipe, param_grid=params, scoring=auc_scorer)
```

Don't overdo it

```
params = dict(  
    feature_selection__k=[2, 3, 4],  
    clf__max_depth=[5, 10, 20],  
    clf__n_estimators=[10, 20, 30]  
)  
grid_search = GridSearchCV(pipe, params, cv=10)
```

$3 \times 3 \times 3 \times 10 = 270$ classifier fits!

Supercharged workflows

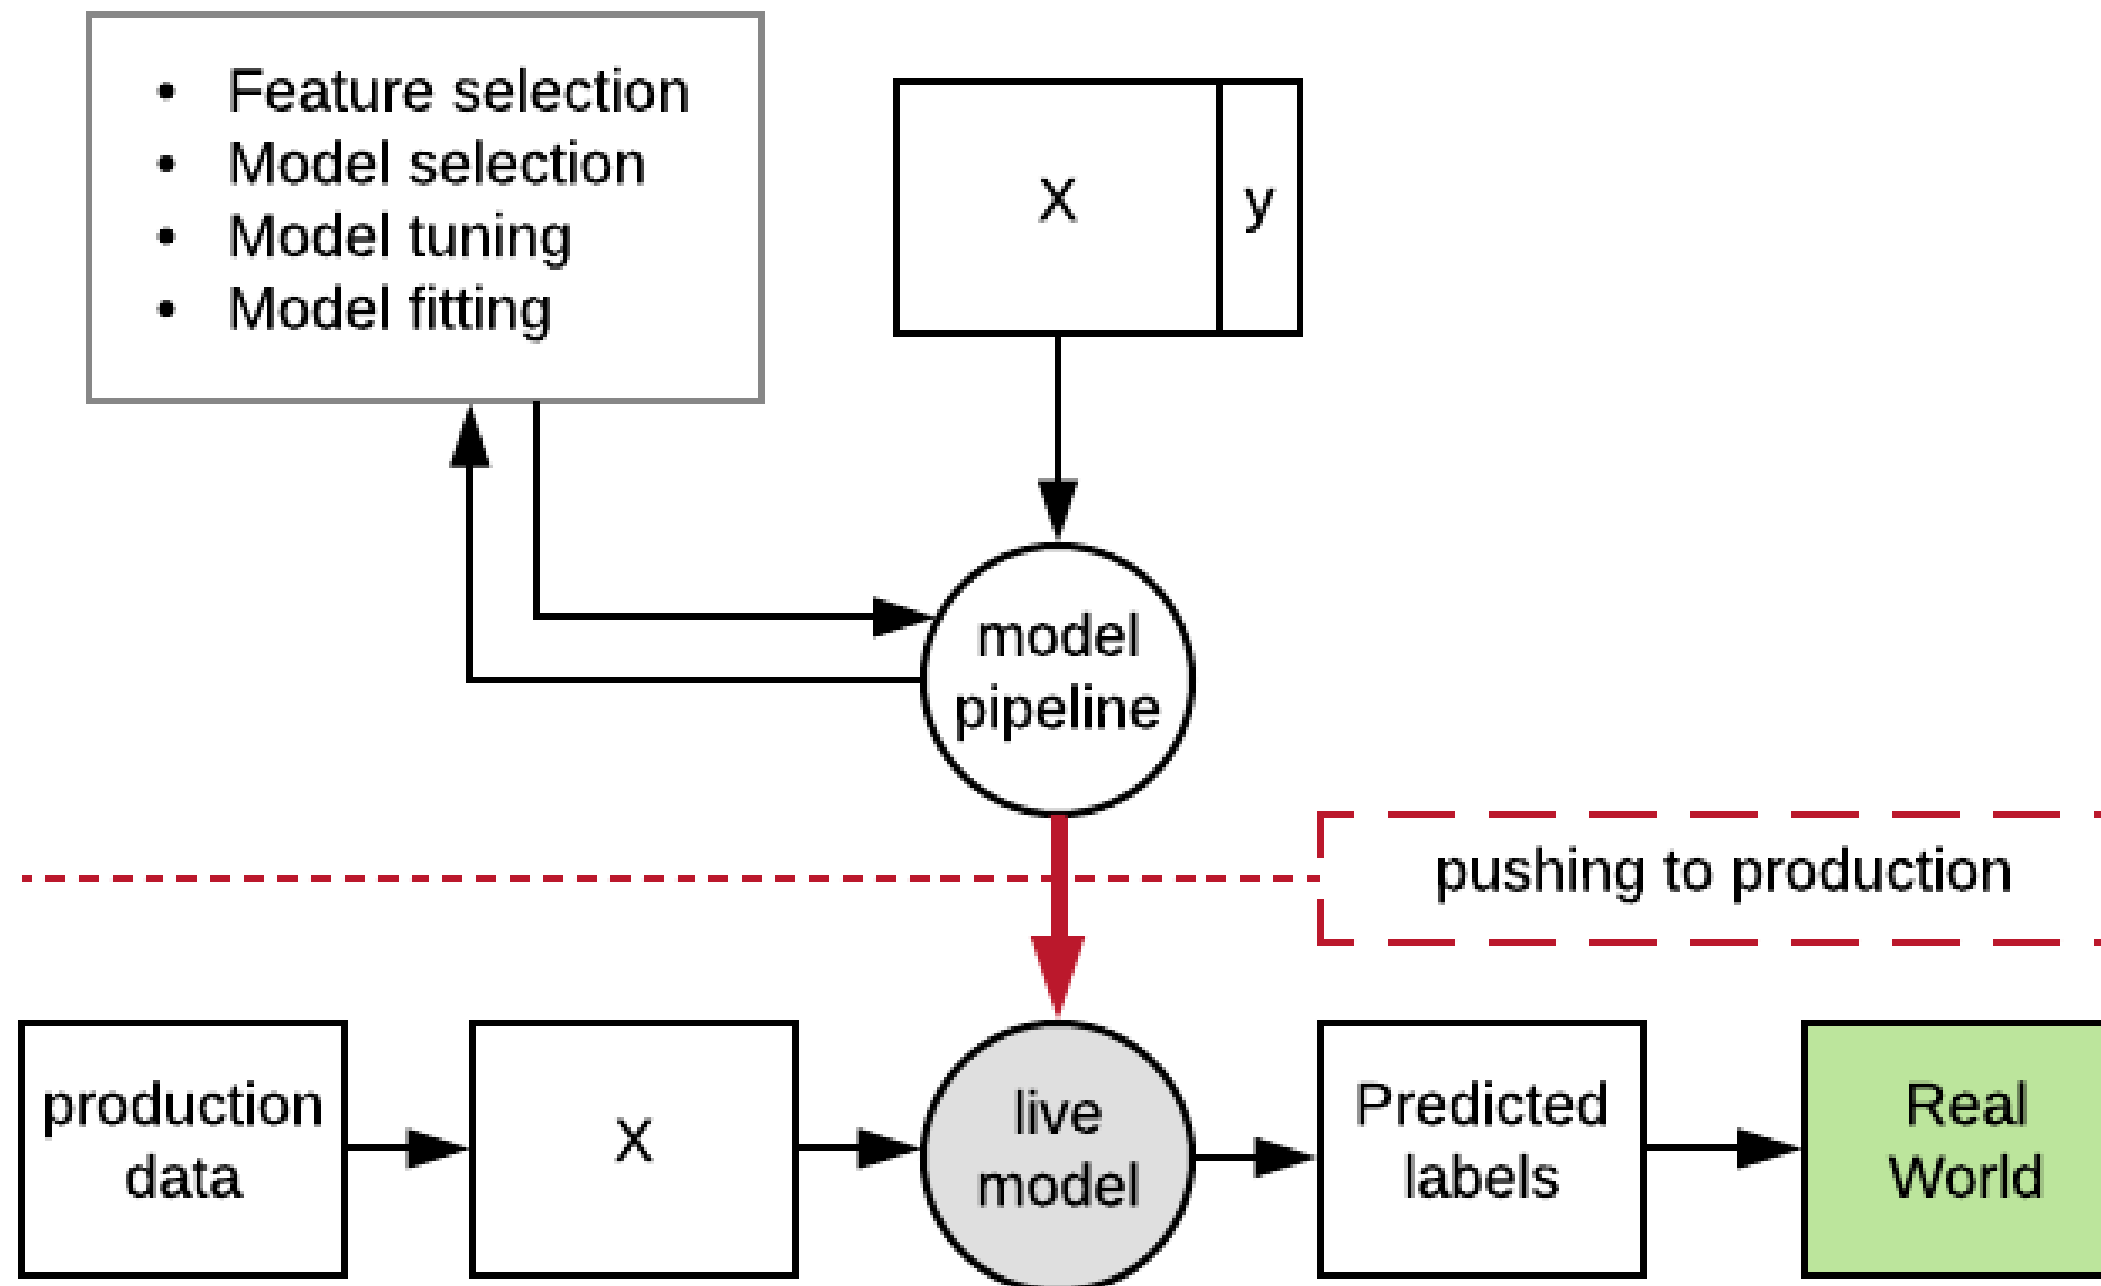
DESIGNING MACHINE LEARNING WORKFLOWS IN PYTHON

Model deployment

DESIGNING MACHINE LEARNING WORKFLOWS IN PYTHON



Dr. Chris Anagnostopoulos
Honorary Associate Professor



Serializing your model

Store a classifier to file:

```
import pickle
clf = RandomForestClassifier().fit(X_train, y_train)
with open('model.pkl', 'wb') as file:
    pickle.dump(clf, file=file)
```

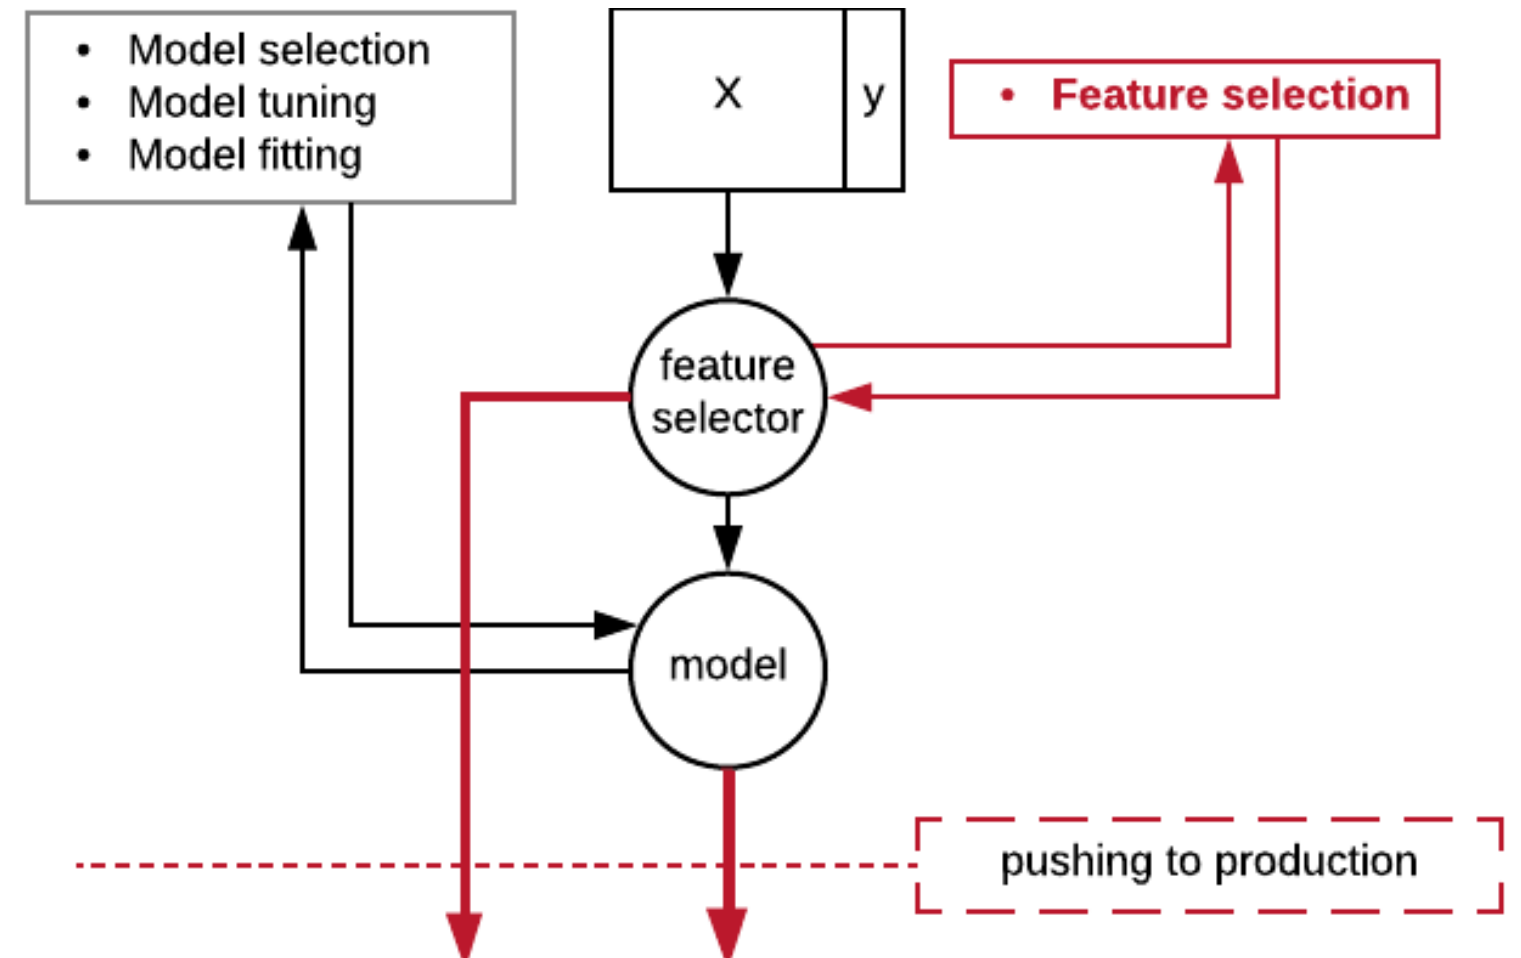
Load it again from file:

```
with open('model.pkl', 'rb') as file:
    clf2 = pickle.load(file)
```

Serializing your pipeline

Development environment:

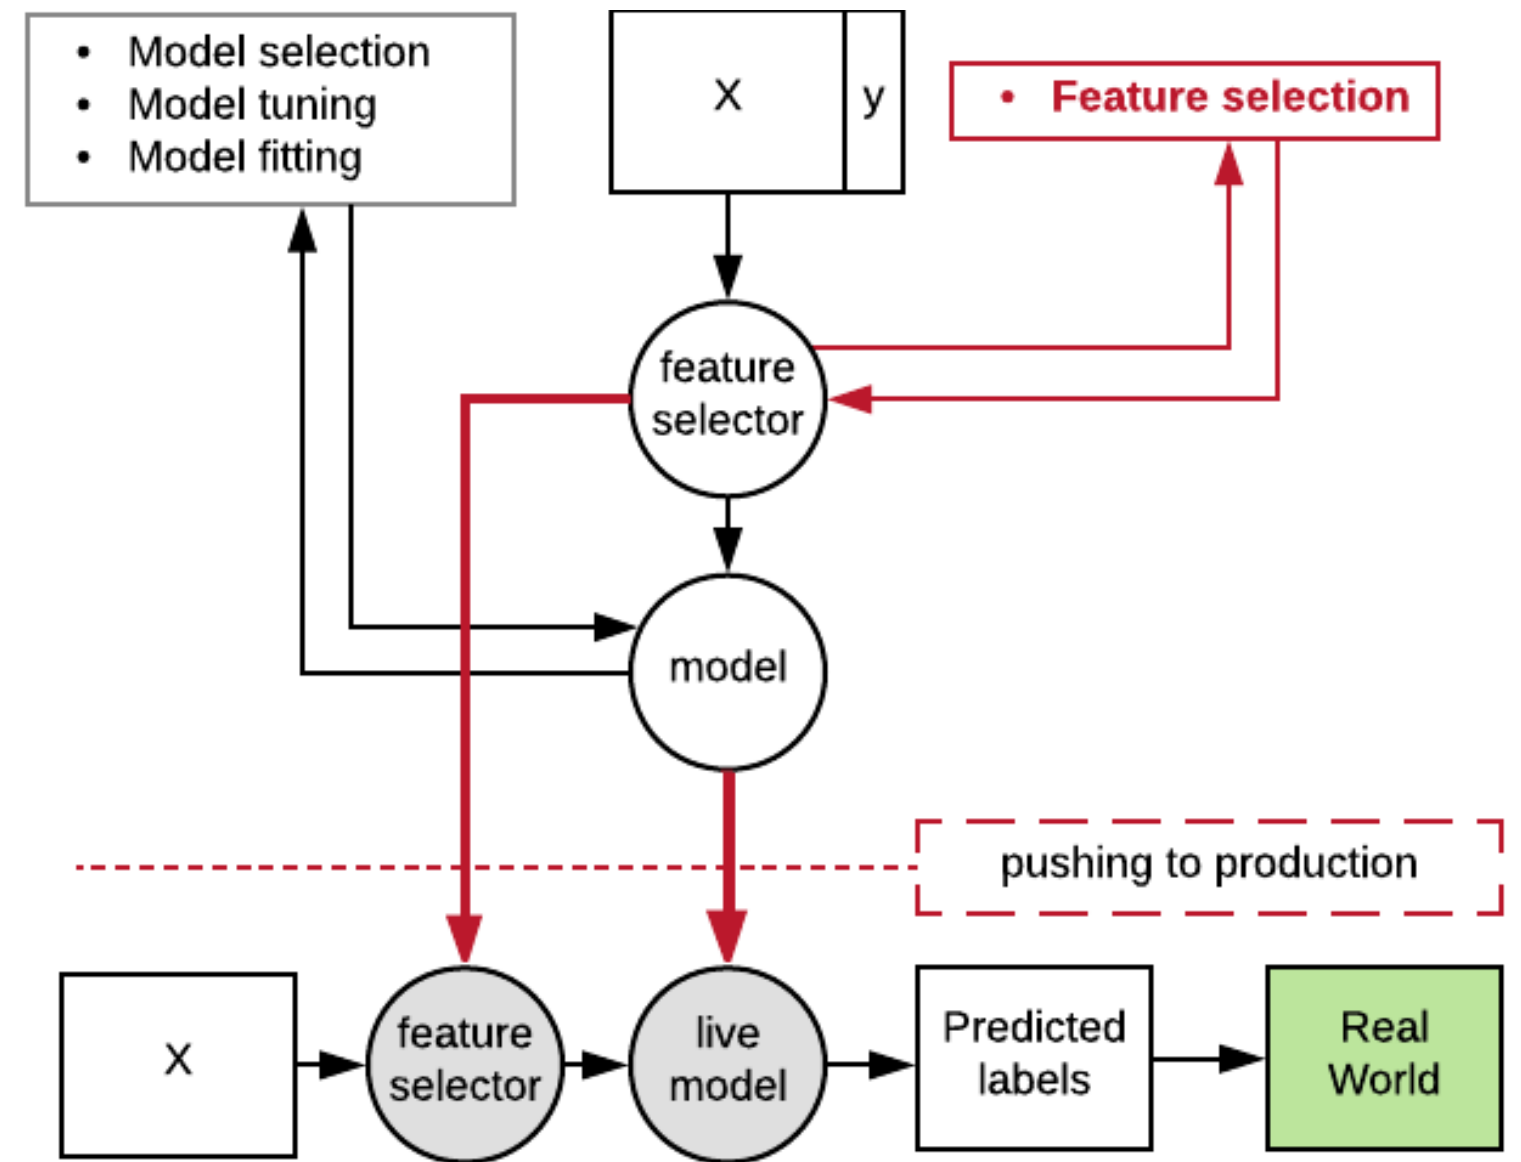
```
vt = SelectKBest(f_classif).fit(  
    X_train, y_train)  
clf = RandomForestClassifier().fit(  
    vt.transform(X_train), y_train)  
with open('vt.pkl', 'wb') as file:  
    pickle.dump(vt)  
with open('clf.pkl', 'wb') as file:  
    pickle.dump(clf)
```



Serializing your pipeline

Production environment:

```
with open('vt.pkl', 'rb') as file:
    vt = pickle.load(vt)
with open('clf.pkl', 'rb') as file:
    clf = pickle.load(clf)
clf.predict(vt.transform(X_new))
```

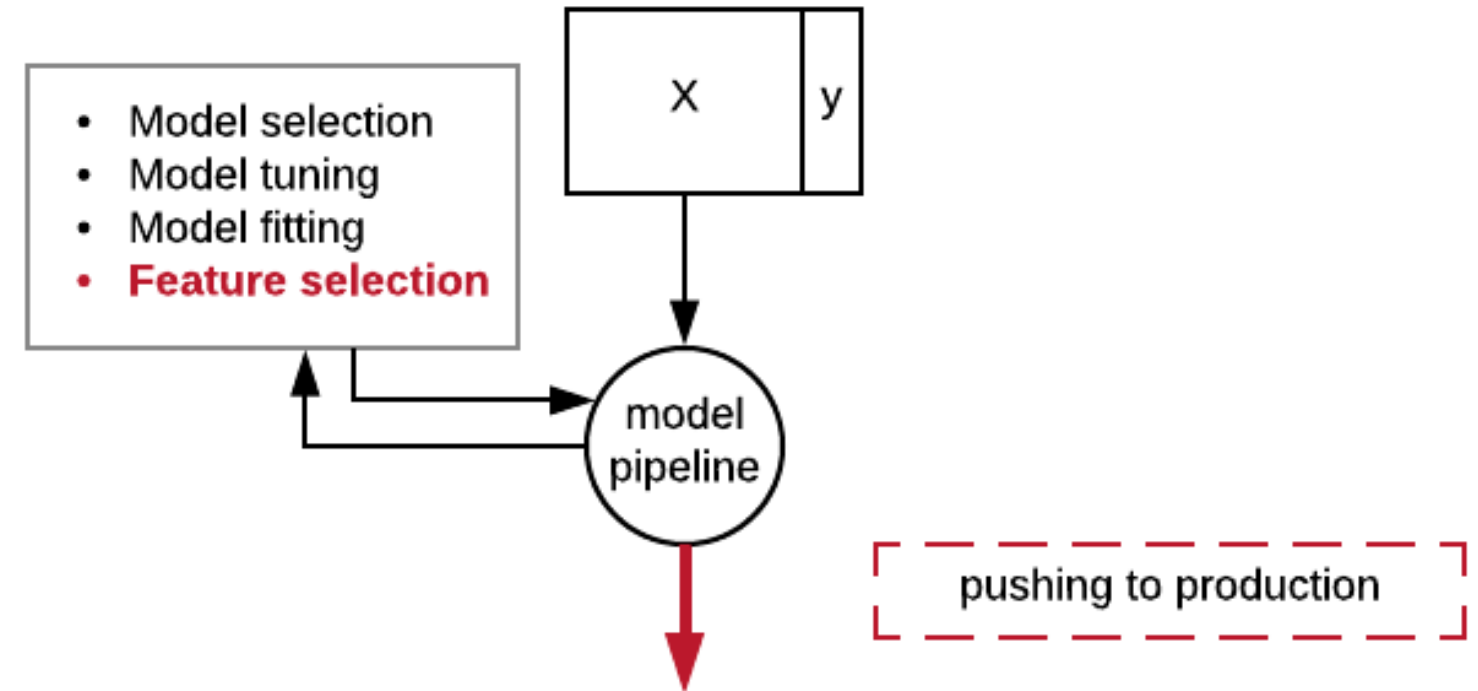


Serializing your pipeline

Development environment:

```
pipe = Pipeline([
    ('fs', SelectKBest(f_classif)),
    ('clf', RandomForestClassifier())
])
params = dict(fs__k=[2, 3, 4],
              clf__max_depth=[5, 10, 20])
gs = GridSearchCV(pipe, params)
gs = gs.fit(X_train, y_train)

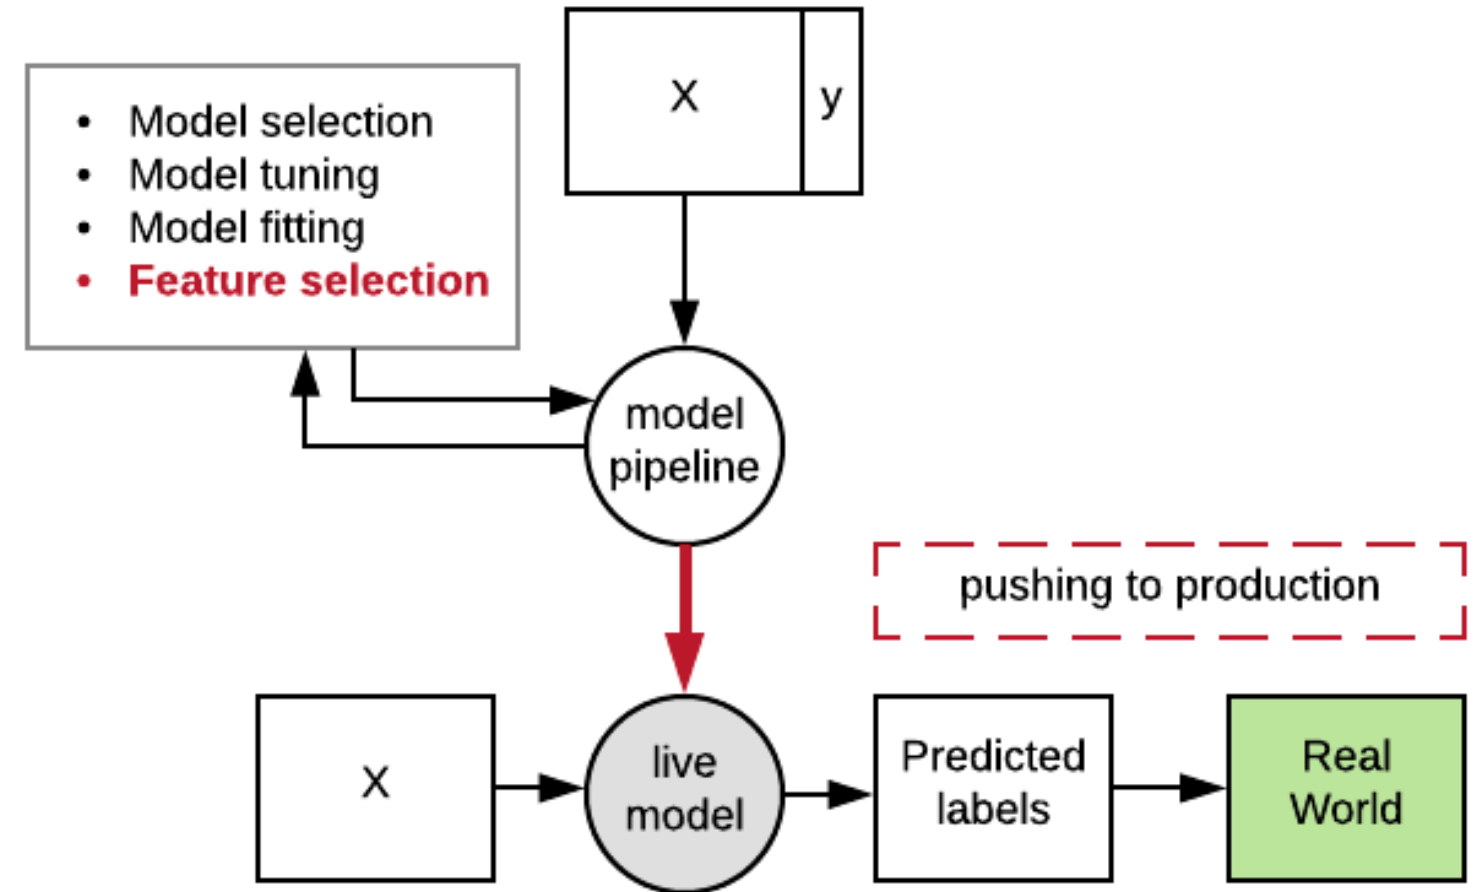
with open('pipe.pkl', 'wb') as file:
    pickle.dump(gs, file)
```



Serializing your pipeline

Production environment:

```
with open('pipe.pkl', 'rb') as file:
    gs = pickle.dump(gs, file)
gs.predict(X_test)
```



Custom feature transformations

	checking_status	duration	...	own_telephone	foreign_worker
0	1	6	...	1	1
1	0	48	...	0	1

```
def negate_second_column(X):  
    Z = X.copy()  
    Z[:, 1] = -Z[:, 1]  
    return Z
```

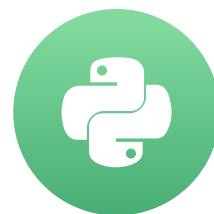
```
pipe = Pipeline([('ft', FunctionTransformer(negate_second_column)),  
                 ('clf', RandomForestClassifier())])
```

Production ready!

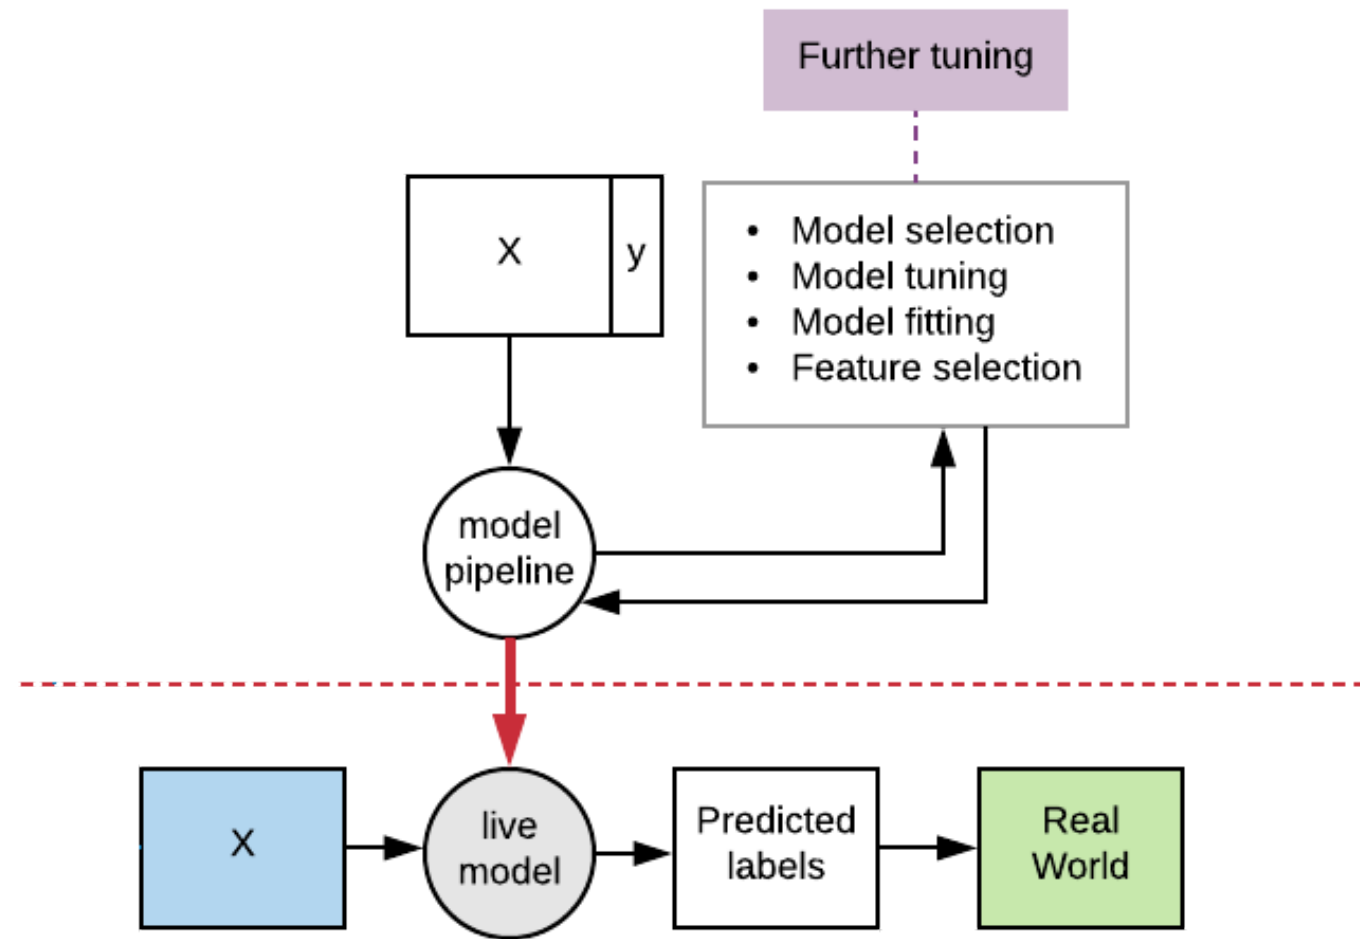
DESIGNING MACHINE LEARNING WORKFLOWS IN PYTHON

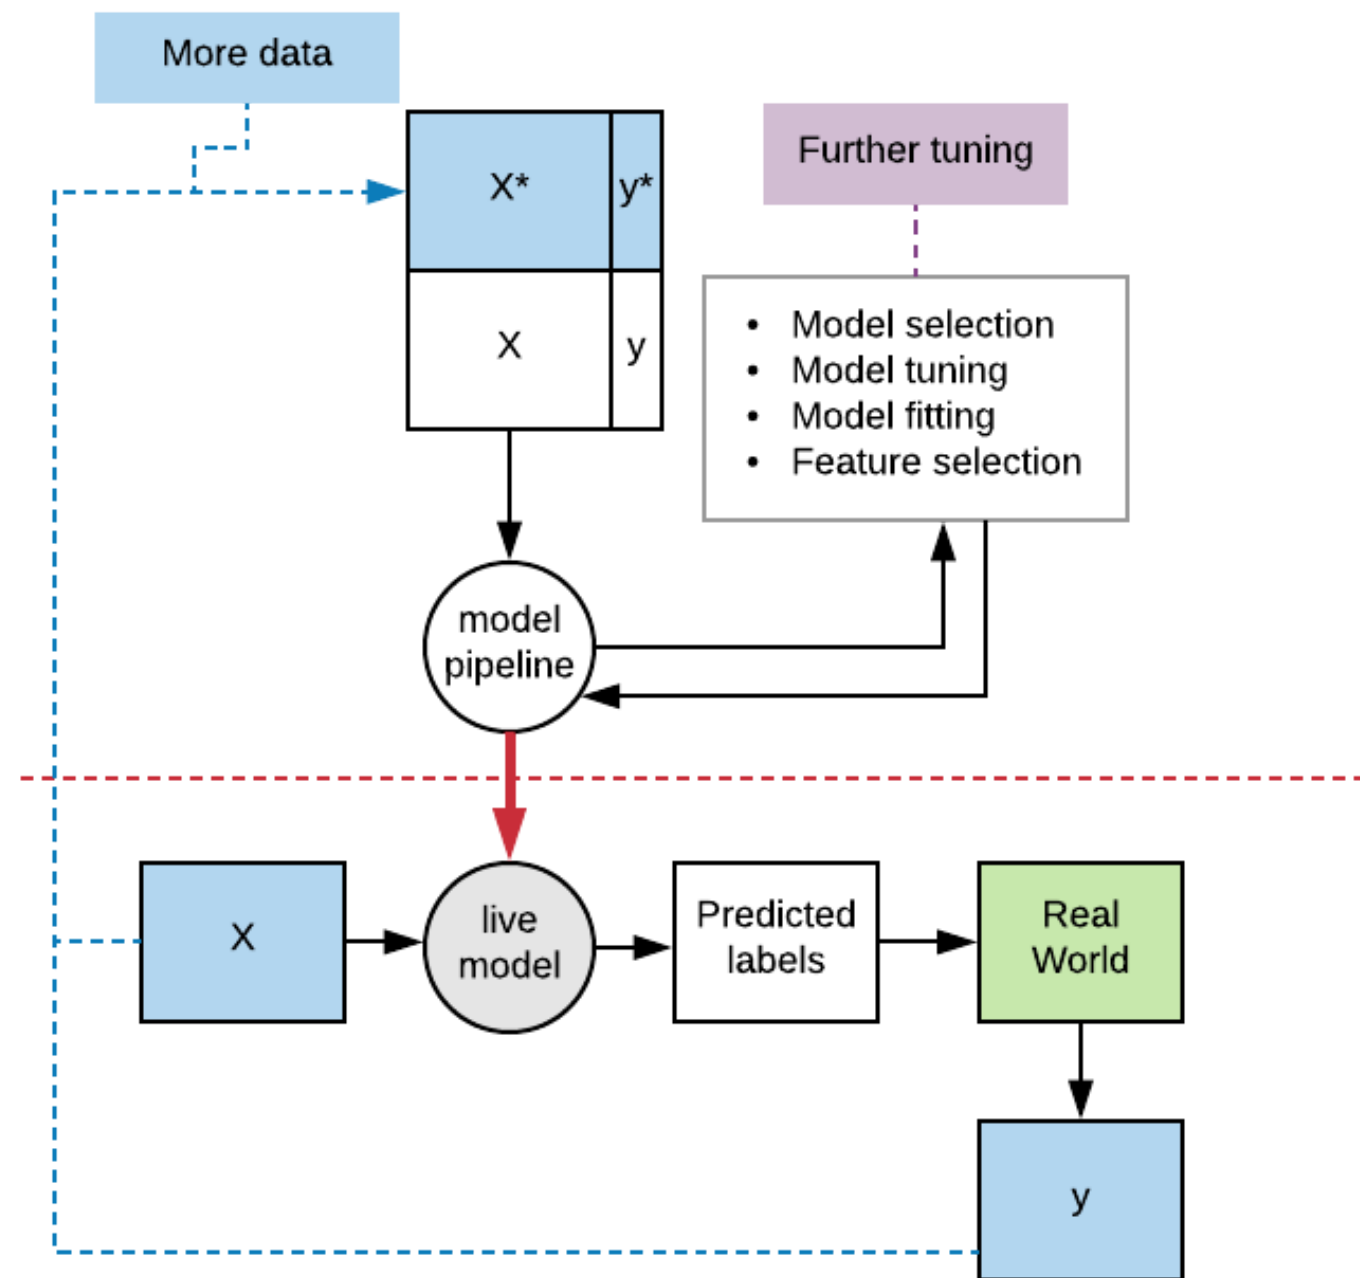
Iterating without overfitting

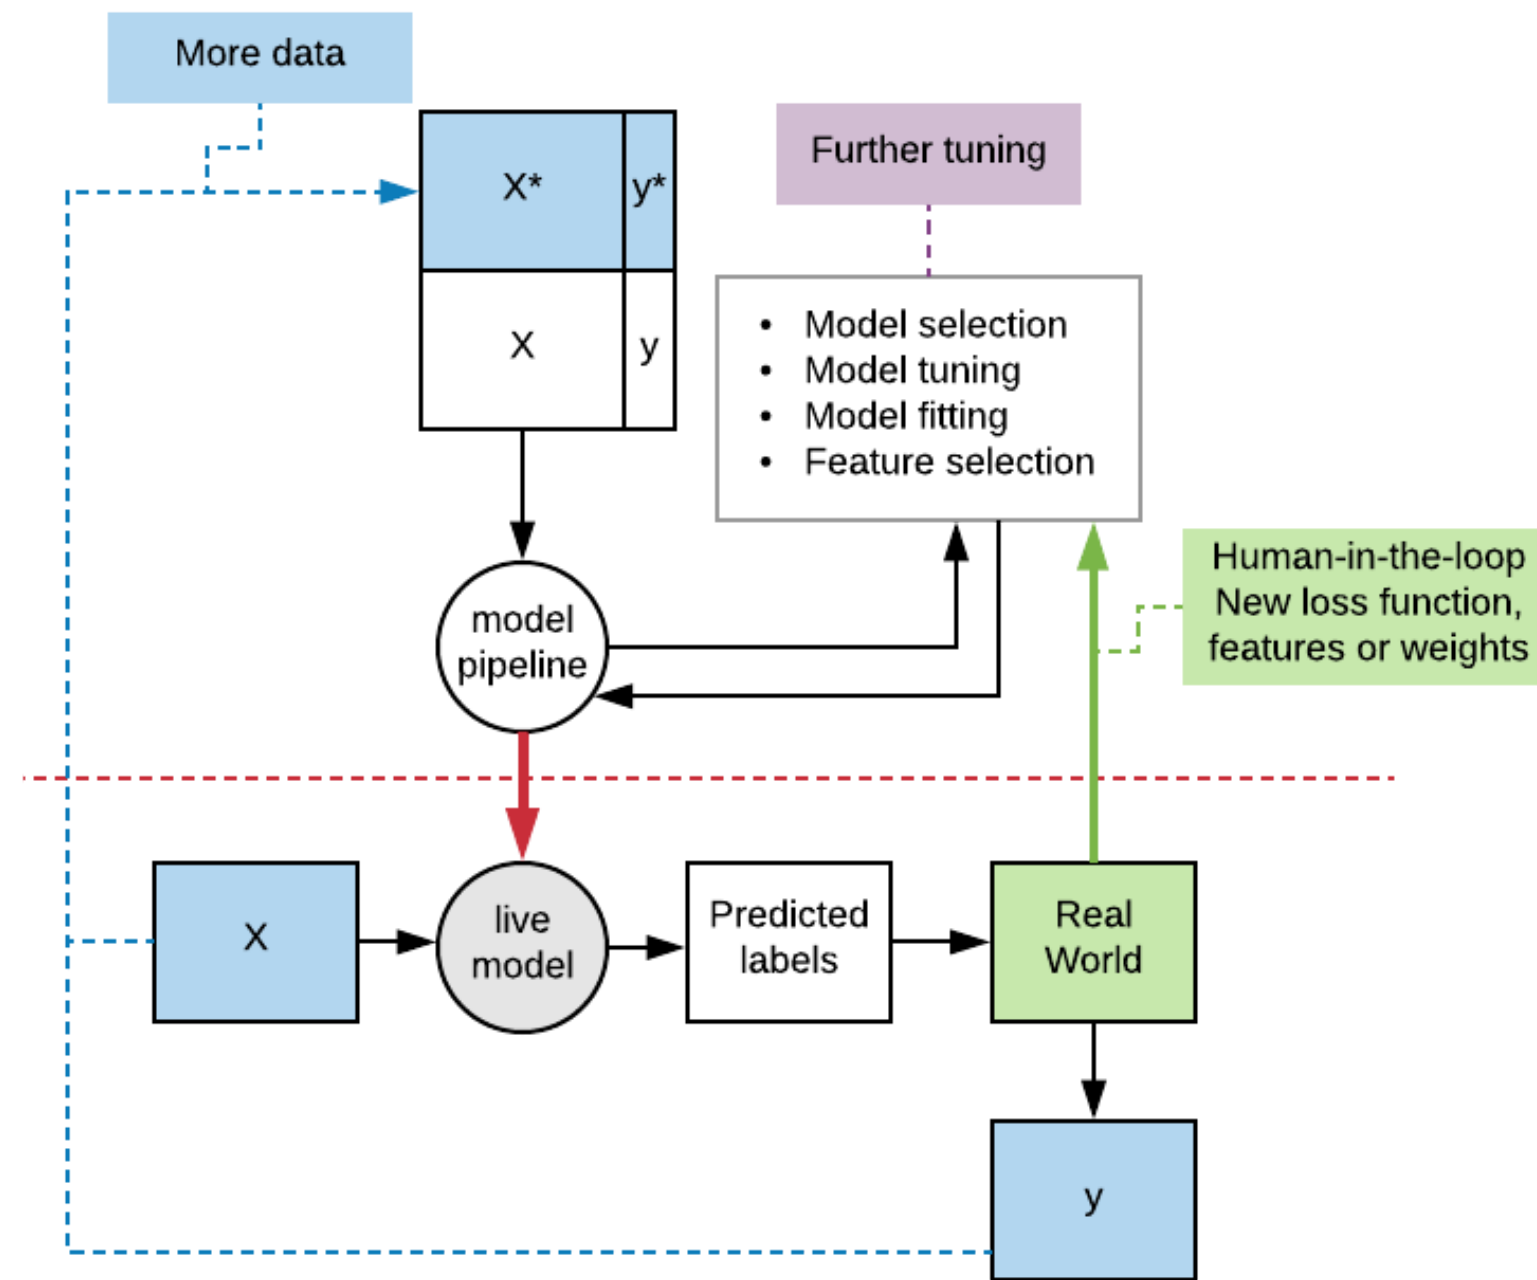
DESIGNING MACHINE LEARNING WORKFLOWS IN PYTHON

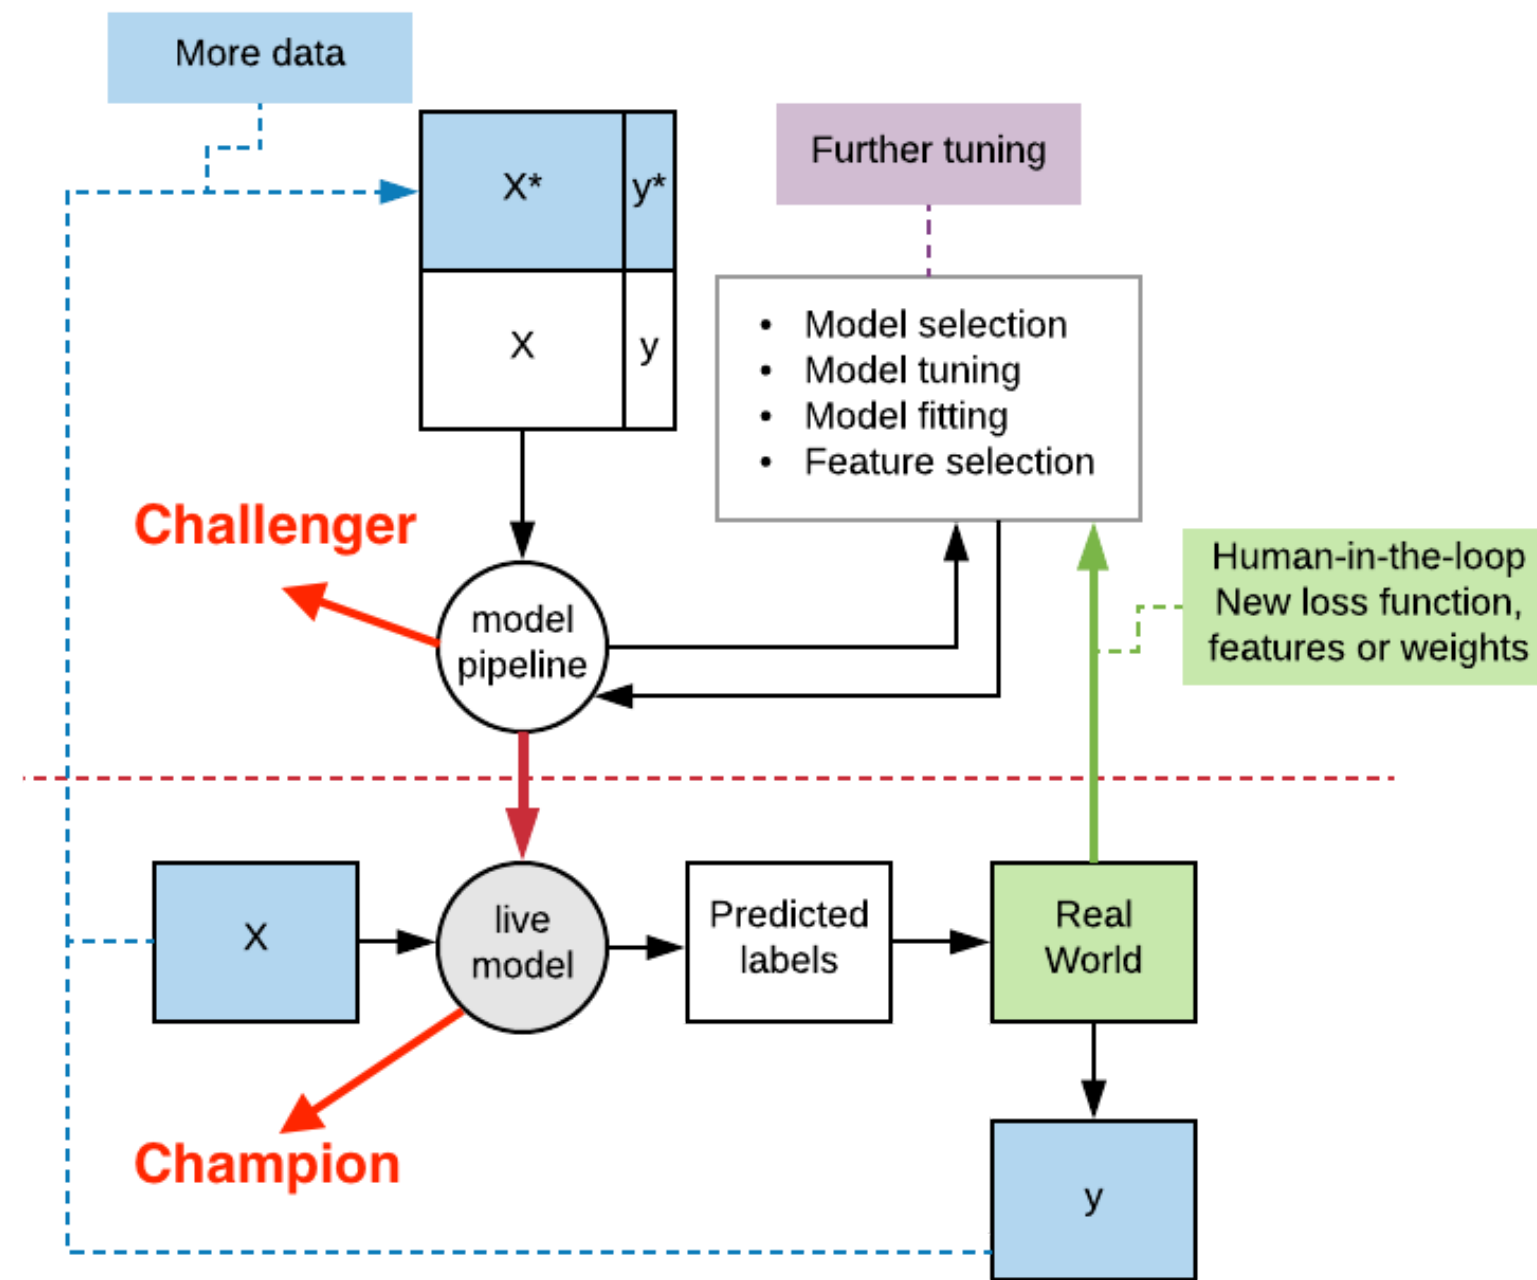


Dr. Chris Anagnostopoulos
Honorary Associate Professor









Cross-validation results

```
grid_search = GridSearchCV(pipe, params, cv=3, return_train_score=True)
gs = grid_search.fit(X_train, y_train)
results = pd.DataFrame(gs.cv_results_)
```

```
results[['mean_train_score', 'std_train_score',
         'mean_test_score', 'std_test_score']]
```


	mean_train_score	std_train_score	mean_test_score	std_test_score
0	0.829	0.006	0.735	0.009
1	0.829	0.006	0.725	0.009
2	0.961	0.008	0.716	0.019
3	0.981	0.005	0.749	0.024
...				


Cross-validation results

	mean_train_score	std_train_score	mean_test_score	std_test_score
0	0.829	0.006	0.735	0.009
1	0.829	0.006	0.725	0.009
2	0.961	0.008	0.716	0.019
3	0.981	0.005	0.749	0.024
4	0.986	0.003	0.728	0.009
5	0.995	0.002	0.751	0.008

Observations:


- Training score much higher than test score.
- The standard deviation of the test score is large.




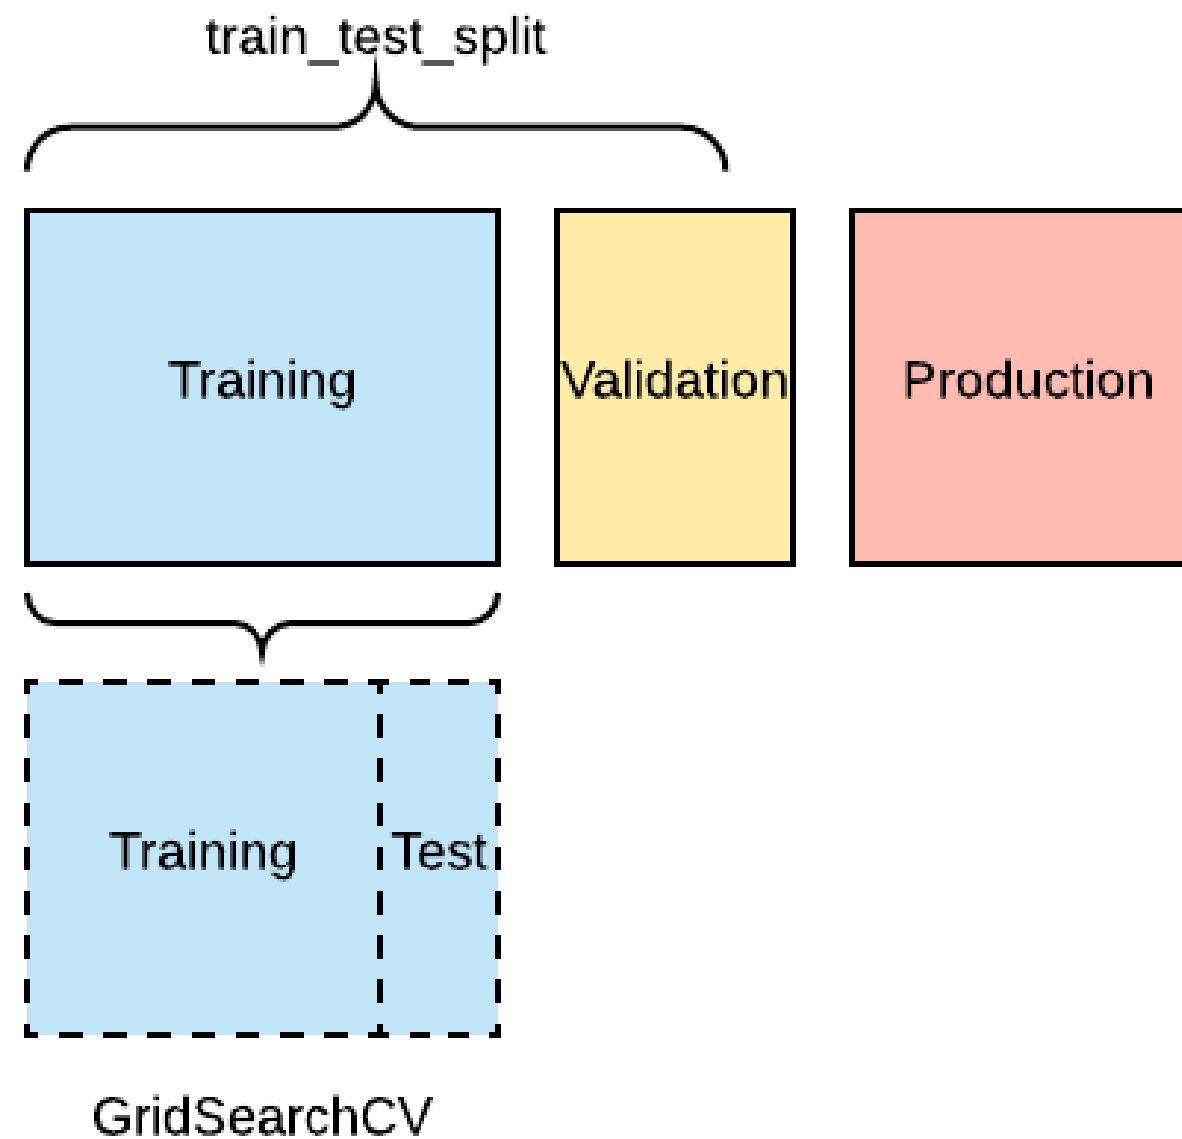


Detecting overfitting

- CV Training Score >> CV Test Score
 - *overfitting in model fitting stage*
 - reduce complexity of classifier
 - get more training data
 - increase cv number
- CV Test Score >> Validation Score
 - *overfitting in model tuning stage*
 - decrease cv number
 - decrease size of parameter grid







"Expert in CV" in your CV!

DESIGNING MACHINE LEARNING WORKFLOWS IN PYTHON

Dataset shift

DESIGNING MACHINE LEARNING WORKFLOWS IN PYTHON



Dr. Chris Anagnostopoulos
Honorary Associate Professor

What is dataset shift?


`elec` dataset:

- 2 years worth of data.
- `class=1` represents price went *up* relative to last 24 hours, and `0` means *down*.


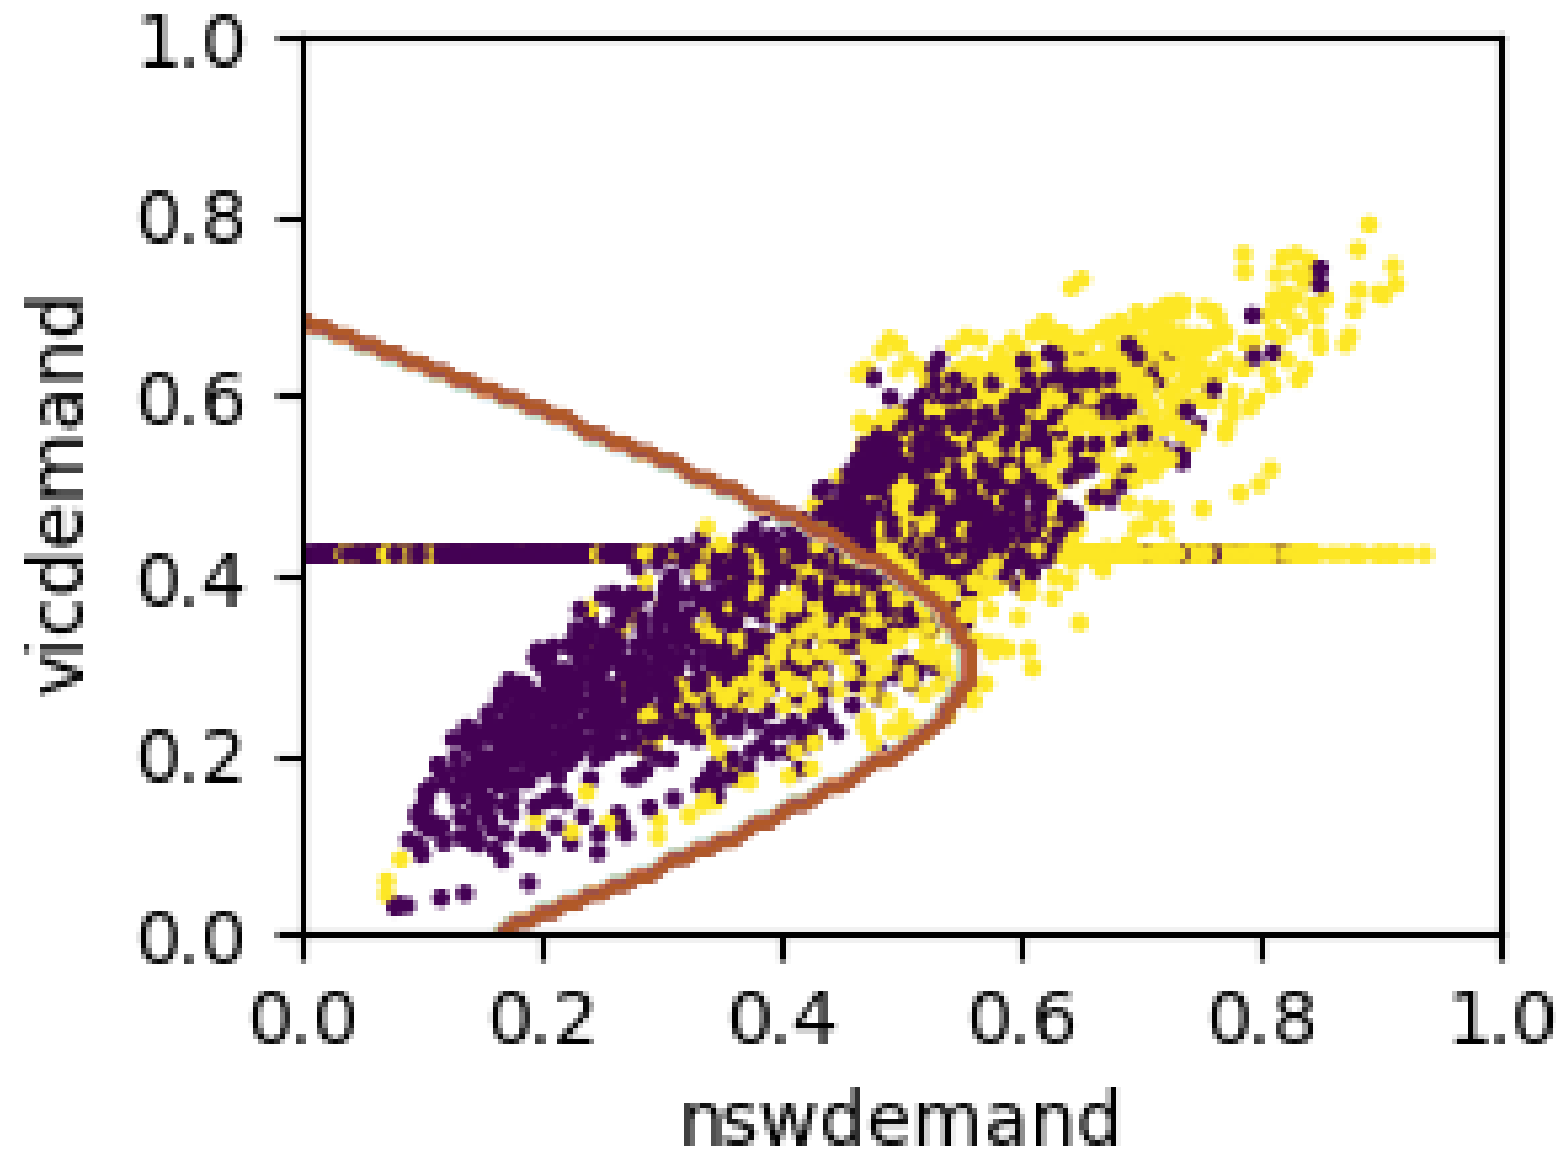
```
   day  period  nswprice  ...  vicedemand  transfer  class
0    2  0.000000  0.056443  ...    0.422915  0.414912     1
1    2  0.553191  0.042482  ...    0.422915  0.414912     0
2    2  0.574468  0.044374  ...    0.422915  0.414912     1
```

```
[3 rows x 8 columns]
```

What is shifting exactly?



What is shifting exactly?



Windows

Sliding window

```
window = (t_now-window_size+1):t_now  
sliding_window = elec.loc>window
```

	day	period	nswprice	...	vicdemand	transfer	class
0	2	0.000000	0.056443	...	0.422915	0.414912	1
1	2	0.553191	0.042482	...	0.422915	0.414912	0
2	2	0.574468	0.044374	...	0.422915	0.414912	1
3	2	0.595745	0.044374	...	0.422915	0.414912	1
4	2	0.617021	0.042482	...	0.422915	0.414912	0
5	2	0.638298	0.040861	...	0.422915	0.414912	0
6	2	0.659574	0.041161	...	0.422915	0.414912	0
7	2	0.680851	0.041161	...	0.422915	0.414912	0
8	2	0.702128	0.041161	...	0.422915	0.414912	0
9	2	0.723404	0.042482	...	0.422915	0.414912	0
10	2	0.765957	0.041041	...	0.422915	0.414912	0
11	2	0.787234	0.040711	...	0.422915	0.414912	0
12	2	0.808511	0.040711	...	0.422915	0.414912	0
13	2	0.829787	0.040861	...	0.422915	0.414912	0
14	2	0.851064	0.041041	...	0.422915	0.414912	0
15	2	0.872340	0.042482	...	0.422915	0.414912	0
16	2	0.893617	0.041161	...	0.422915	0.414912	0
17	2	0.914894	0.051489	...	0.422915	0.414912	1
18	2	0.936170	0.056443	...	0.422915	0.414912	1
19	2	0.957447	0.054642	...	0.422915	0.414912	1

Expanding window

```
window = 0:t_now  
expanding_window = elec.loc>window
```

	day	period	nswprice	...	vicdemand	transfer	class
0	2	0.000000	0.056443	...	0.422915	0.414912	1
1	2	0.553191	0.042482	...	0.422915	0.414912	0
2	2	0.574468	0.044374	...	0.422915	0.414912	1
3	2	0.595745	0.044374	...	0.422915	0.414912	1
4	2	0.617021	0.042482	...	0.422915	0.414912	0
5	2	0.638298	0.040861	...	0.422915	0.414912	0
6	2	0.659574	0.041161	...	0.422915	0.414912	0
7	2	0.680851	0.041161	...	0.422915	0.414912	0
8	2	0.702128	0.041161	...	0.422915	0.414912	0
9	2	0.723404	0.042482	...	0.422915	0.414912	0
10	2	0.765957	0.041041	...	0.422915	0.414912	0
11	2	0.787234	0.040711	...	0.422915	0.414912	0
12	2	0.808511	0.040711	...	0.422915	0.414912	0
13	2	0.829787	0.040861	...	0.422915	0.414912	0
14	2	0.851064	0.041041	...	0.422915	0.414912	0
15	2	0.872340	0.042482	...	0.422915	0.414912	0
16	2	0.893617	0.041161	...	0.422915	0.414912	0
17	2	0.914894	0.051489	...	0.422915	0.414912	1
18	2	0.936170	0.056443	...	0.422915	0.414912	1
19	2	0.957447	0.054642	...	0.422915	0.414912	1

Dataset shift detection

```
# t_now = 40000, window_size = 20000
clf_full = RandomForestClassifier().fit(X, y)
clf_sliding = RandomForestClassifier().fit(sliding_X, sliding_y)
```

```
# Use future data as test
test = elec.loc[t_now:elec.shape[0]]
test_X = test.drop('class', 1); test_y = test['class']
```


```
roc_auc_score(test_y, clf_full.predict(test_X))
roc_auc_score(test_y, clf_sliding.predict(test_X))
```

```
0.775
```

```
0.780
```

Window size

```
for w_size in range(10, 100, 10):  
    sliding = arrh.loc[  
        (t_now - w_size + 1):t_now  
    ]  
    X = sliding.drop('class', 1)  
    y = sliding['class']  
    clf = GaussianNB()  
    clf.fit(X, y)  
    preds = clf.predict(test_X)  
    roc_auc_score(test_y, preds)
```



Domain shift

arrhythmia dataset:

	age	sex	height	...	chV6_TwaveAmp	chV6_QRSA	chV6_QRSTA	class
0	75	0	190	...	2.9	23.3	49.4	0
1	56	1	165	...	2.1	20.4	38.8	0
2	54	0	172	...	3.4	12.3	49.0	0
3	55	0	175	...	2.6	34.6	61.6	1
4	75	0	190	...	3.9	25.4	62.8	0

[5 rows x 280 columns]

More data is not always better!

DESIGNING MACHINE LEARNING WORKFLOWS IN PYTHON